

## Python Gradaute - Parsing SwissProt

“Without requirements or design, programming is the art of adding bugs to an empty text file.” - Louis Srygley

Create a Python program called “swisstake.py” that processes a SwissProt-formatted file as a positional argument. It should have a *required* `-k|--keyword` argument of the keyword to match in the “keyword” field of the input record in order to determine which sequences to “take” (hence the name). It should also have an *optional* `-s|--skip` argument to “skip” records with given taxa (which could be many so `nargs='+'`), as well as an *optional* `-o|--output` argument to where to write the output in FASTA format (default “out.fa”).

If the given input file is not a file, it should die with “XXX” is not a file’.

```
$ ./swisstake.py -h
usage: swisstake.py [-h] [-s STR [STR ...]] -k STR [-o FILE] FILE
```

Filter Swissprot file for keywords, taxa

positional arguments:

FILE Uniprot file

optional arguments:

-h, --help show this help message and exit

-s STR [STR ...], --skip STR [STR ...]  
Skip taxa (default: )

-k STR, --keyword STR  
Take on keyword (default: None)

-o FILE, --output FILE  
Output filename (default: out.fa)

```
$ ./swisstake.py swiss.txt
```

```
usage: swisstake.py [-h] [-s STR [STR ...]] -k STR [-o FILE] FILE
```

```
swisstake.py: error: the following arguments are required: -k/--keyword
```

```
$ ./swisstake.py -k proteome foo
```

```
"foo" is not a file
```

```
$ ./swisstake.py swiss.txt -k "complete proteome" -s Metazoa FUNGI viridiplantae
Processing "swiss.txt"
```

```
Done, skipped 14 and took 1. See output in "out.fa".
```

```
$ ./swisstake.py swiss.txt -k "complete proteome" -s metazoa fungi
Processing "swiss.txt"
```

```
Done, skipped 13 and took 2. See output in "out.fa".
```

## BioPython SwissProt Record

A FASTA record had three attributes: ID, description, and sequence. A SwissProt record has considerably more which will make sense once you look at the file.

There are at least two ways I've found to parse a SwissProt record. One is use `SeqIO.parse(fh, 'swiss')` which gives you a record very similar to a FASTA record which has an `annotations` attribute which is a dictionary that looks like this:

```
>>> import pprint
>>> pp = pprint.PrettyPrinter(indent=4)
>>> pp.pprint(rec.annotations)
{ 'accessions': ['P13813'],
  'date': '01-JAN-1990',
  'date_last_annotation_update': '20-JAN-2016',
  'date_last_sequence_update': '01-JAN-1990',
  'entry_version': 42,
  'keywords': ['Malaria', 'Repeat'],
  'ncbi_taxid': ['5850'],
  'organism': 'Plasmodium knowlesi',
  'protein_existence': 2,
  'references': [ Reference(title='Cloning and characterization of an abundant Plasmodium
  'sequence_version': 1,
  'taxonomy': [ 'Eukaryota',
                'Alveolata',
                'Apicomplexa',
                'Aconoidasida',
                'Haemosporida',
                'Plasmodiidae',
                'Plasmodium',
                'Plasmodium (Plasmodium)']}]}
```

The other way is use the `Bio.SwissProt` module which has attributes for the same kind of information though sometimes called slightly different names, e.g.:

```
>>> sw1.organism_classification
['Eukaryota', 'Alveolata', 'Apicomplexa', 'Aconoidasida', 'Haemosporida', 'Plasmodiidae', 'P
```

Cf:

- [https://biopython.readthedocs.io/en/latest/Tutorial/chapter\\_uniprot.html](https://biopython.readthedocs.io/en/latest/Tutorial/chapter_uniprot.html)
- <http://biopython.org/DIST/docs/api/Bio.SwissProt.Record-class.html>

However you choose to parse, you should be able to pass the tests. FWIW, I used the first method.

## Sets

We’ve talked about dictionaries quite a bit, and for this exercise I think you’ll want something that is a natural extension of a dictionary called `set()` which is just a dictionary where the values are all 1. If you have two lists, you can test for equality:

```
>>> a = ['foo', 'bar']
>>> b = ['foo', 'bar']
>>> a == b
True
>>> c = ['bar', 'foo']
>>> a == c
False
```

The list `c` has the same members but in a different order, so the lists definitely are not the same; however, if all you cared about what if the two lists shared the same items, you could sort them:

```
>>> sorted(a) == sorted(c)
True
```

But what if you wanted to know if there was some overlap. Clearly you can’t use equality:

```
>>> d = ['foo', 'bar', 'baz']
>>> a == d
False
```

You have to individually check each element of `a` to see if they are in `d`:

```
>>> [e for e in a if e in d]
['foo', 'bar']
>>> [e for e in d if e in a]
['foo', 'bar']
>>> any([e for e in a if e in d])
True
```

That is the “intersection” of the two lists. The “difference” would be:

```
>>> [e for e in a if e not in d]
[]
>>> [e for e in d if e not in a]
['baz']
```

The “union” would be everything in both lists, which we can’t easily do in one line of code; however, if we convert these lists to sets, then we can do all those calculations easily:

```
>>> a = set(['foo', 'bar'])
>>> d = set(['foo', 'bar', 'baz'])
```

```
>>> a.union(d)
set(['baz', 'foo', 'bar'])
>>> a.intersection(d)
set(['foo', 'bar'])
>>> a.difference(d)
set([])
>>> d.difference(a)
set(['baz'])
```

Keep this in mind when you are trying to find if there is an intersection of the taxa you are given with the taxa that are in the record.

## Test Suite

A passing test suite looks like this:

```
$ make test
python3 -m pytest -v test.py
===== test session starts =====
platform darwin -- Python 3.6.8, pytest-4.2.0, py-1.7.0, pluggy-0.8.1 -- /anaconda3/bin/python
cachedir: .pytest_cache
rootdir: /Users/kyclark/work/worked_examples/07-grad-swissprot, inifile:
plugins: remotedata-0.3.1, openfiles-0.3.2, doctestplus-0.2.0, arraydiff-0.3
collected 3 items

test.py::test_usage PASSED [ 33%]
test.py::test_bad_input PASSED [ 66%]
test.py::test_good_input1 PASSED [100%]

===== 3 passed in 1.75 seconds =====
```