

Unix Basics

Most of bioinformatics happens on Unix-like operating systems. Almost all our tools run from the Unix command line, so bioinformaticians must know how to move data around, run programs, and chain the output of one program into another to create analysis pipelines.

NB: When you see a **\$** given in the example prompts, it is a metacharacter indicating that this is the prompt for a normal (not super-user) account. You should type/copy/paste all the stuff *after* the **\$**. If you ever see a prompt with “**#**” in a tutorial, it’s indicating a command that should be run as the super-user/root account, e.g., installing some software into a system-wide directory so it can be shared by all users.

Common Unix Commands

I assume you are on a command line by now, so let’s look at some commands.

- **whoami**: reports your username
- **w**: shows who is currently on a system
- **man**: show the manual page for a command
- **echo**: say something
- **cowsay**: have a cow say something
- **env**: print your environment
- **printenv**: print some/all of your environment
- **which/type**: tells you the location of a program
- **touch**: create an empty regular file
- **file**: briefly describe a file or directory
- **pwd**: print working directory, where you are right now
- **ls**: list files in current directory
- **find**: find files or directories matching some conditions
- **locate**: find files using a database, requires daemon to run
- **cd**: change directory (with no arguments, **cd \$HOME**)
- **cp**: copy a file (or “**cp -r**” to copy a directory)
- **mv**: move a file or directory
- **mkdir**: create a directory
- **rmdir**: remove a directory
- **rm**: remove a file (or “**rm -r**” to remove a directory)
- **cat**: concatenate files (cf. <http://porkmail.org/era/unix/award.html>)
- **column**: arrange text into columns
- **paste**: merge lines of files
- **sort**: sort text or numbers
- **uniq**: remove duplicates *from a sorted list*
- **sed**: stream editor for altering text
- **awk/gawk**: pattern scanning and processing language

- **grep**: global regular expression program (maybe?), cf. https://en.wikipedia.org/wiki/Regular_expression
- **history**: look at past commands, cf. CTRL-R for searching your history directly from the command line
- **head**: view the first few (10) lines of a file
- **tail**: view the last (10) lines of a file
- **comm**: find lines in common/unique in two sorted files
- **top**: view the programs taking the most system resources (memory, I/O, time, CPUs, etc.), cf. “htop”
- **ps**: view the currently running processes
- **cut**: select columns from the output of a program
- **wc**: (character) word (and line) count
- **more/less**: pager programs that show you a “page” of text at a time; cf. <https://unix.stackexchange.com/questions/81129/what-are-the-differences-between-most-more-and-less/81131>
- **bc**: calculator
- **df**: report file system disk space usage; useful to find a place to land your data
- **du**: report disk usage; recommend “du -shc”; useful to identify large directories that need to be removed
- **ssh**: secure shell, like telnet only with encryption
- **scp**: secure copy a file from/to remote systems using ssh
- **rsync**: remote sync; uses scp but only copies data that is different
- **ftp**: use “file transfer protocol” to retrieve large data sets (better than HTTP/browsers, esp for getting data onto remote systems)
- **ncftp**: more modern FTP client that automatically handles anonymous logins
- **wget**: web get a file from an HTTP location, cf. “wget is not a crime” and Aaron Schwartz
- **|**: pipe the output of a command into another command
- **>, >>**: redirect the output of a command into a file; the file will be created if it does not exist; the single arrow indicates that you wish to overwrite any existing file, while the double-arrows say to append to an existing file
- **<**: redirect contents of a file into a command
- **nano**: a very simple text editor; until you’re ready to commit to vim or emacs, start here
- **md5sum**: calculate the MD5 checksum of a file
- **diff**: find the differences between two files

Chaining commands

Most Unix commands use STDIN (“standard in”), STDOUT (“standard out”), and STDERR (“standard error”). For instance, the **env** program will show you

key/value pairs that describe your environment – things like your user name (\$USER), your shell (\$SHELL), your current working directory (\$PWD). It can be quite a long list, so you could send the output (STDOUT) of `env` to `head` to see just the first few lines:

```
$ env | head
TERM_PROGRAM=iTerm.app
TERM=xterm-256color
SHELL=/bin/bash
TMPDIR=/var/folders/0h/vjzky052qx4p70trn2p2h400000gn/T/
PERL5LIB=/Users/kyclark/work/imicrobe/lib
Apple_PubSub_Socket_Render=/private/tmp/com.apple.launchd.MoI0Cra0uS/Render
TERM_PROGRAM_VERSION=3.2.6
OLDPWD=/Users/kyclark/work/biosys-analytics/lectures
TERM_SESSION_ID=w2t0p0:34E89330-9CEE-4FA6-9039-3674CBDE4655
USER=kyclark
```

If you provide a file name as an argument to `head`, it will work on that:

```
$ head /usr/share/dict/words
A
a
aa
aal
aalii
aam
Aani
aardvark
aardwolf
Aaron
```

Without any arguments, `head` assumes you must want it to read from STDIN, which is why you can chain any program's STDOUT into `head`. For instance, you could use `grep` to look for lines with the word "TERM" in them:

```
$ env | grep TERM
TERM_PROGRAM=iTerm.app
TERM=xterm-256color
TERM_PROGRAM_VERSION=3.2.6
TERM_SESSION_ID=w2t0p0:34E89330-9CEE-4FA6-9039-3674CBDE4655
ITERM_PROFILE=Default
ITERM_SESSION_ID=w2t0p0:34E89330-9CEE-4FA6-9039-3674CBDE4655
COLORTERM=truecolor
```

If you are fortunate enough to have the `fortune` and `cowsay` programs on your system, you can do this:

```
$ fortune | cowsay
```

```

/ Somebody ought to cross ball point pens \
| with coat hangers so that the pens will |
\ multiply instead of disappear.           /
-----
      \  ^__^
       \  (oo)\_______
          (__)\       )\/\
              ||----w |
              ||     ||

```

Manual pages

The `man` program will show you the manual page for a program, if it exists. Just type `man <program>`, e.g., `man wget`. Inside a manpage, you can use the `/` to search for a string. Use `q` to “quit” `man`. Most programs will also show you a help/usage document if you run them with `-h`, `--help`, or `-help`. I often find it useful to `grep` the help, e.g.:

```

$ wget --help | grep clobber
  -nc, --no-clobber          skip downloads that would download to
  --unlink                  remove file before clobber

```

Pronunciation

- `/`: “slash”; the thing leaning the other way is a “backslash”
- `sh`: “shuh” or “ess-ach”
- `etc`: “et-see”
- `usr`: “user”
- `src`: “source”
- `#`: “hash” (NOT “hashtag”) or “pound”
- `$`: “dollar”
- `!`: “bang”
- `#!`: “shebang”
- `^`: “caret”
- `PID`: “pid” (not pee-eye-dee)
- `~`: “twiddle” or “tilde”; shortcut to your home directory when alone, shortcut to another user’s home directory when used like “~bhurwitz”

Variables

You will see things like `$USER` and `$HOME` that start with the `$` sign. These are variables because they can change from person to person, system to system. On most systems, my username is “kyclark” but I might be “kclark” or “kyclark1” on others, but on all systems `$USER` refers to whatever value is defined for my username. Similarly, my `$HOME` directory might be “/Users/kyclark,” “/home1/03137/kyclark,” or “/home/u20/kyclark,” but I can always refer to the idea of my home directory with the variable `$HOME`.

Control sequences

If you launch a program that won’t stop, you can use CTRL-C to send an “interrupt” signal to the program. If it is well-behaved, it should stop, but it may not. You can open another terminal on the machine and run `ps -fu $USER` to find all the programs you are running.

```
$ ps -fu $USER
UID          PID  PPID  C  STIME TTY          TIME CMD
kyclark    31718 31692  0  12:16 ?           00:00:00 sshd: kyclark@pts/75
kyclark    31723 31718  0  12:16 pts/75      00:00:00 -bash
kyclark    33265 33247  0  12:16 ?           00:00:00 sshd: kyclark@pts/86
kyclark    33277 33265  1  12:16 pts/86      00:00:00 -bash
kyclark    33792 33277  9  12:17 pts/86      00:00:00 vim run.p6
kyclark    33806 31723  0  12:17 pts/75      00:00:00 ps -fu kyclark
```

The PID is the “process ID” and the PPID is the “parent process ID.” In the above table, let’s assume my “vim” session was locked up. I could `kill 33792`. If in a reasonable amount of time (a minute or so) that doesn’t work, common wisdom says you `kill -9` to really, really tell it to shut down, but:

No no no. Don’t use `kill -9`.

It doesn’t give the process a chance to cleanly:

- shut down socket connections
- clean up temp files
- inform its children that it is going away
- reset its terminal characteristics

and so on and so on and so on.

Generally, send 15, and wait a second or two, and if that doesn’t work, send 2, and if that doesn’t work, send 1. If that doesn’t, REMOVE THE BINARY because the program is badly behaved!

Don’t use `kill -9`. Don’t bring out the combine harvester just to tidy up the flower pot. cf. <http://porkmail.org/era/unix/award.html#>

uuk9letter

Along with CTRL-C, you should learn about CTRL-Z to put a process into the background. This could be handy if, say, you were in an editor, you make a change to your code, you CTRL-Z to background the editor, you run the script to see if it worked, then you **fg** to bring it back to the foreground or **bg** it to have it resume running in the background. I would consider this a sub-optimal work environment, but it's fine if you were for some reason limited to a single terminal window.

Generally if you want a program to run in the background, you would launch it from the command line with an ampersand (“&”) at the end:

```
$ my-background-prog.sh &
```

Lastly, know that most Unix programs interpret CTRL-D as the end-of-input signal. You can use this to send the “exit” command to most any interactive program, even your shell. Here's a way to enter some text into a file directly from the command line without using a text editor. After typing the last line (i.e., type “chickens.<Enter>”), type CTRL-D:

```
$ cat > wheelbarrow
so much depends
upon
```

```
a red wheel
barrow
```

```
glazed with rain
water
```

```
beside the white
chickens.
```

```
<CTRL-D>
```

```
$ cat wheelbarrow
so much depends
upon
```

```
a red wheel
barrow
```

```
glazed with rain
water
```

```
beside the white
chickens.
```

Handy command line shortcuts

- Tab: hit the Tab key for command completion; hit it early and often!
- **!!**: (bang-bang) execute the last command again
- **!\$**: (bang-dollar) the last argument from your previous command line (think of the \$ as the right anchor in a regex)
- **!^**: (bang-caret) the first argument from your previous command line (think of the ^ as the left anchor in a regex)
- CTRL-R: reverse search of your history
- Up/down cursor keys: go backwards/forwards in your history
- CTRL-A, CTRL-E: jump to the start, end of the command line when in emacs mode (default)

N.B.: If you are on a Mac, it's easy to remap your (useless) CAP-SLOCK key to CTRL. Much less strain on your hand as you will find you need CTRL quite a bit, even more so if you choose emacs for your \$EDITOR.

Altering your \$PATH

As you see above, “env” will list all the key-value pairs defining your environment. For instance, everyone has a \$HOME directory that you can see with `echo $HOME`. The exact location of \$HOME can vary among systems, e.g.:

- Mac: /Users/kyclark
- Ocelot: /home/u20/kyclark
- Stampede: /home1/03137/kyclark

Your \$PATH setting is extremely important as it defines the directory locations that will be searched (in order) to find programs. Here's my \$PATH on the UA HPC:

```
[hpc:login3@~]$ echo $PATH | sed "s:/:/:\\n/g"
/rsgroups/bhurwitz/hurwitzlab/bin:
/home/u20/kyclark/.cargo/bin:
/home/u20/kyclark/.local/bin:
/cm/local/apps/gcc/6.1.0/bin:
/cm/shared/uaapps/pbspro/18.2.1/sbin:
/cm/shared/uaapps/pbspro/18.2.1/bin:
/opt/TurboVNC/bin:
/cm/shared/uabin:
/usr/lib64/qt-3.3/bin:
/cm/local/apps/environment-modules/4.0.0/bin:
/usr/local/bin:
/bin:
/usr/bin:
```

```
/usr/local/sbin:
/usr/sbin:
/sbin:
/sbin:
/usr/sbin:
/cm/local/apps/environment-modules/4.0.0/bin
```

I’ve used “sed” to add a newline after each colon so you can more easily see that the directories are separated by colons. Notice that I have the shared “hurwitzlab” directory first in my path. Much of our work will require access to tools that are not installed by default on the HPC. You could build them into your own `$HOME` directory, but it will be easier if you just add this shared directory to your `$PATH`. From the command line, you can do this:

```
export PATH=/rsgrps/bhurwitz/hurwitzlab/bin:$PATH
```

You just told your shell (bash) to set the `$PATH` variable to our “hurwitzlab” directory and then whatever it was set to before. Obviously we want this to happen each time we log in, so we can add this command to `$HOME/.bashrc`:

```
echo "export PATH=/rsgrps/bhurwitz/hurwitzlab/bin:$PATH" >> ~/.bashrc
```

N.B., “dotfiles” are files with names that begin with a dot. They are normally hidden from view unless you use `ls -a` to list “all” files. A single dot `.` means the current directory, and two dots `..` mean the parent directory. Your “`.bashrc`” (or maybe “`.profile`” or maybe “`.bash_profile`” depending on your system) file is read every time you login to your system, so you can remember your customizations. “`Rc`” may mean “resource configuration,” but who really knows?

As you find or create useful programs that you would like to have available globally on your system (i.e., not just in the current working directory), you can create a location like `$HOME/bin` (or my preferred `$HOME/.local/bin`) and add this to your `$PATH` as well. You can add as many directories as you like (within reason).

Aliases

Sometimes you’ll find you’re using a particular command quite often and want to create a shortcut. You can assign any command to a single “alias” like so:

```
alias cx='chmod +x'
alias up2='cd ../../'
alias up3='cd ../../../../'
```

If you execute this on the command line, the alias will be saved until you log out. Adding these lines to your `.bashrc` will make it available every time you log in. When you make a change and want the shell to bring those into the

current environment, you need to **source** the file. The command `.` is an alias for **source**:

```
$ source ~/.bashrc
$ . ~/.bashrc
```

Permissions

When you execute `ls -l`, you'll see the "long" listing of the contents of a directory similar to this:

```
-rwxr-xr-x  1 kycklark  staff      174 Aug  9 20:21 abs.py*
drwxr-xr-x 14 kycklark  staff     476 Aug  3 12:14 anaconda3/
```

The first column of data contains a wealth of information represent in 10 bits. The first bit is:

- "-" for a regular file
- "d" for a directory
- "l" for a symlink (like a shortcut)

The other nine bits are broken into sets of three bits that represent the permissions for the "user," "group," and "other." The "abs.py" is a regular file we can tell from the first dash. The next three bits show "rwx" which means that the user ("kycklark") has read, write, and execute permissions for this file. The next three bits show "r-x" meaning that the group ("staff") can read and execute the file only. The same is true for all others.

When you create a file, the normal default is that it is not executable. You must specifically tell Unix to change the mode of the file using the "chmod" command. Often it's enough to say:

```
$ chmod +x myprog.sh
```

To turn on the execute bits for everyone, but it's possible to have much finer control of the permissions. If you only want user and group to have execute, then do:

```
$ chmod ug+x myprog.sh
```

Removing is done with a "-", so any combination of `[ugo] [+ -] [rwx]` will usually get you what you want.

Sometimes you may see instructions to `chmod 775` a file. This is using octal notation where the three bits "rwx" correspond to the digits "421," so the first "7" is "4+2+1" which equals "rwx" whereas the "5" = "4+1" so only "rw":

user			group			other				
r	w	x	r	w	x	r	w	x		
4	2	1		4	2	1		4	2	1

```

+ + +   + + +   + - +
= 7     = 7     = 5

```

Therefore “chmod 775” is the same as:

```

$ chmod -rwx myfile
$ chmod ug+rwx myfile
$ chmod o+rw myfile

```

When you create ssh keys or config files, you are instructed to `chmod 600`:

```

user    group    other
r w x   r w x   r w x
4 2 1 | 4 2 1 | 4 2 1
+ + -   - - -   - - -
= 6     = 0     = 0

```

Which means that only you can read or write the file, and no one else can do anything with it. So you can see that it can be much faster to use the octal notation.

When you are trying to share data with your colleagues who are on the same system, you may put something into a shared location but they complain that they cannot read it or nothing is there. The problem is most likely permissions. The “uask” setting on a system determines the default permissions, and it may be that the directory and/or files are readable only by you. It may also be that you are not in a common group that you can use to grant permission, in which case you can either:

- politely ask your sysadmin to create a new group OR
- `chmod 777` the directory, which is probably the worst option as it makes the directory completely accessible to anyone to do anything. In short, don’t do this unless you really don’t care if someone accidentally or maliciously wipes out your data.

File system layout

The top level of a Unix file system is “/” which is called “root.” Confusingly, there is also an account named “root” which is basically the super-user/sysadmin (systems administrator). Unix has always been a multi-tenant system ...

Installing software

Much of the time, “bioinformatics” seems like little more than installing software and chaining them together with scripts. Sometimes you may be lucky enough

to have a “sysadmin” (systems administrator) who can assist you, but most of the time you’ll find yourself needing to take care of business yourself.

My suggestions for installing software are (in order):

Sysadmin

Go introduce yourself to your sysadmins. Take them to lunch or order them some pizza or drop off some good beer or whiskey. Whatever it takes to be on good terms because a good sysadmin who is responsive to your needs is an enormous help. If they are willing to install software for you, that is the way to go. Often, though, this is a task far beneath them, and they would expect you to be able to fend for yourself. They may provide `sudo` (<https://xkcd.com/149/>) privileges to allow you to install software into shared locations (e.g., `/usr/local`), but it’s more likely they would expect you to install into your `$HOME`.

Package managers

There are several package management systems for Linux and OSX including `apt-get`, `yum`, `homebrew`, `macports`, and more. These usually relieve the problems of software compatibility and shared libraries. Unless you have `sudo` to install globally, you can configure to install into your `$HOME`.

Binary installations

Quite often you’ll be happy to find that the maintainers of the software you need have gone to the trouble to build binary distributions for your system, which is likely to be a generic 64-bit Linux platform. Often you can just download the binaries and put them into your `$PATH`. There is usually a “README” or “INSTALL” file that will explain exactly what to do. To use the binaries, you can:

- 1) Always refer to the full path to the binary
- 2) Place them into a directory in your `$PATH` like `$HOME/.local/bin`
- 3) Add the new directory to your `$PATH`

Source installations

Installing from source usually means downloading a “tarball” (“tar” = “tape archive,” a container of files, that is then compressed with a program like “gzip” to create a “.tar.gz” or “.tgz” file extension), running “configure” to figure out how it can build on your system, and then “make” to build the binaries. Usually

you will run “make install” to put the binaries into their proper directory, but sometimes you just “make” and copy the files yourself.

The basic steps for installing into your `$HOME` are usually:

```
$ tar xvf package.tgz
$ ./configure --prefix=$HOME/.local
$ make && make install
```

When I’m in an environment with a directory I can share with my team (like the UA HPC), I’ll configure the package to install into that shared space so that others can use the program. When I’m on a system like “stampede” where I cannot share with others, I’ll usually install into my `$HOME/.local` or some sort of “work” directory.